

DoC 437 - 2008

Distributed Algorithms

Part 5: Asynchronous Algorithms

Asynchronous models

- Asynchronous models make no assumptions about *time*

execution time

no bound on time to execute a local process step

communication time

no bound on message transmission delay

no synchronized clocks

no rounds to structure the algorithm

Asynchronous models

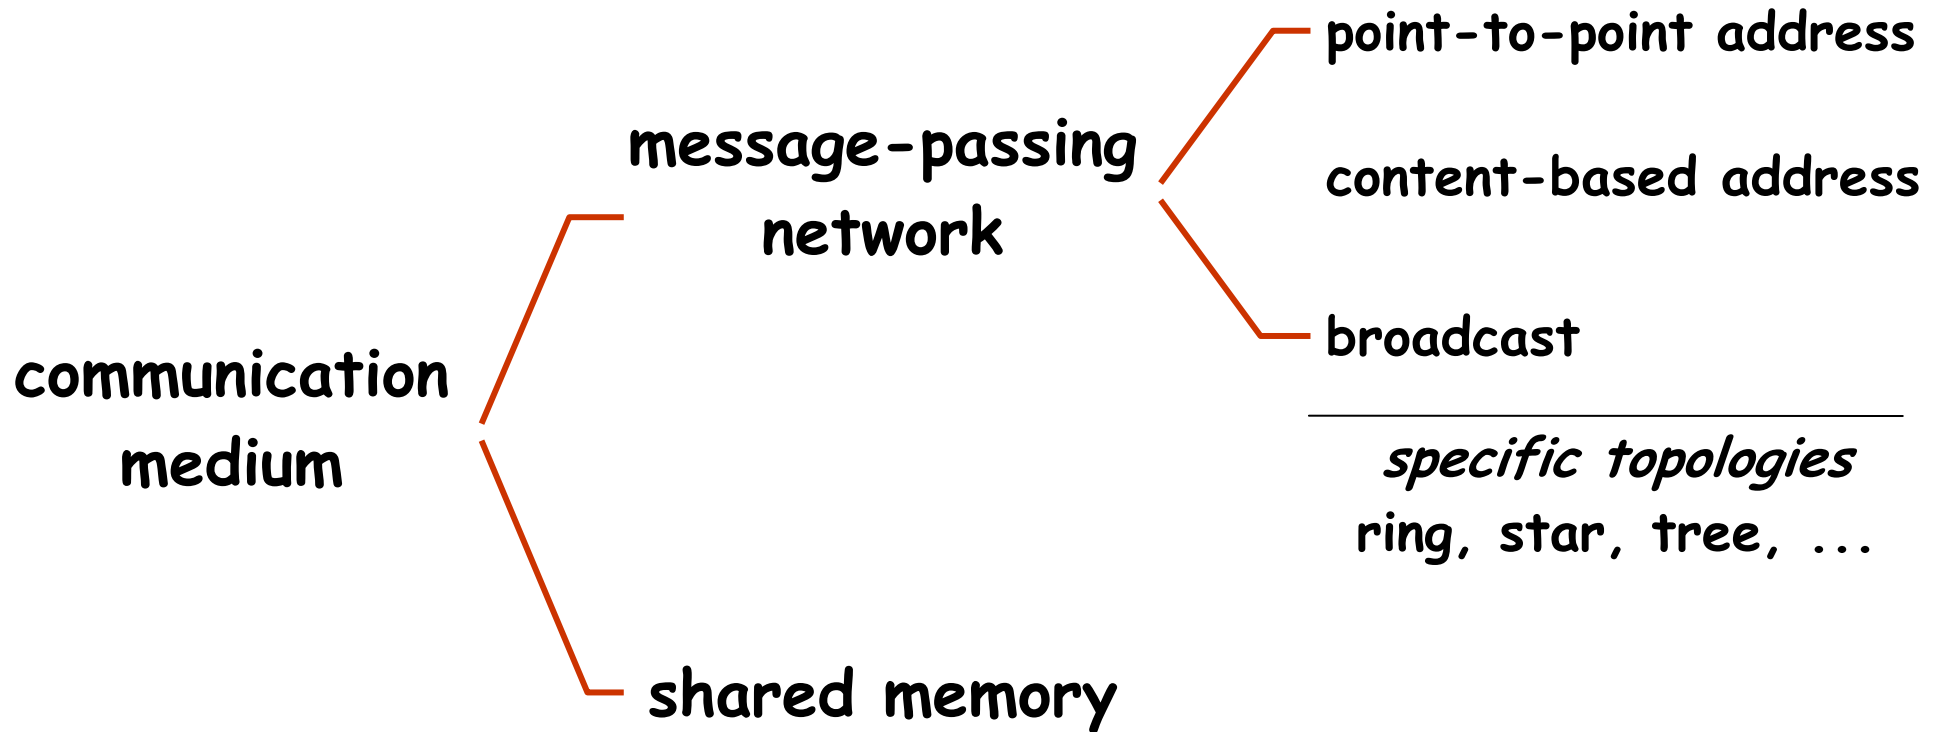
- Main differences from synchronous models involve *liveness conditions*

uncertainty caused by asynchrony and distribution
more general distributed algorithms, as they embody
weaker assumptions


more general than most distributed systems

usually more difficult to program than under the
synchronous model

Recall: communication medium




Asynchronous shared memory



- Dual of asynchronous network model
can freely translate between them
- A simpler model about which to reason
- Further simplification
ignore faults

Shared variables in shared memory



- Each shared variable is defined by:


a set of *values*

an *initial* value

read and/or write

capabilities available
to processes

Process interaction



- Each process is defined by:

a set of *states*

a *start* state

a state and variable
transition function

that is, an algorithm
represented as a *state
machine*

Mutual exclusion for two processes

- The basic model

using asynchronous shared memory, two processes, p_1 and p_2 , compete for access to their critical regions (i.e., access to shared resource)

- The desired outcome (properties)

mutual exclusion (*safety*)

freedom from deadlock (*safety*)

freedom from starvation (*liveness*)

Mutual exclusion, informally

- When process p_i wants to enter its critical region \mathcal{C} , it sets a flag f_i
- If the other process's flag is set, p_i does not enter \mathcal{C} , but instead resets f_i and tries again
- If the other process's flag is not set, p_i enters \mathcal{C}
- p_i resets f_i after leaving \mathcal{C}

Code skeleton

shared variables: $f(i:1..2)$: Boolean, initially false, writable by p_i , readable by all

process p_i :

try:


$f(i) := \text{true}$

 if ($f((i\%2)+1)$) then {

$f(i) := \text{false}$; go to try}

 else { enter/use/exit critical region }

$f(i) := \text{false}$



Mutual exclusion and deadlock

- What is an example of an error in the algorithm that would lead to...

violation of deadlock freedom?

violation of mutual exclusion?

Safety property: mutual exclusion

- Need to argue that it is not possible for p_1 and p_2 to both be in critical region C at the same time
- Proof sketch (by contradiction)

assume p_1 and p_2 both in C , so f_1 and f_2 must be set and remain set from before C to E

if p_1 set flag first...

p_2 must remain in T and cannot enter C (same for p_2 and p_1)
hence, contradiction

if p_1 and p_2 set flags "simultaneously"...

then both remain in T and neither can enter C
hence, contradiction

What about starvation?

- Recall: freedom from starvation is a liveness property
any process that reaches T eventually enters \mathcal{C}
- Recall: processes modeled as state machines
we might assume "fair choice"
if a choice over a set of transitions is executed infinitely often, then every transition in the set is executed infinitely often

Modeling "fairness"

- *Strong fairness* (a form of fair choice)

for every transition, if it is enabled infinitely often, it is taken infinitely often

- *Weak fairness*

for every transition, if it is enabled continuously from some point on, it is taken infinitely often

- *No fairness*

for every transition, even if enabled continuously from some point, it may not be taken unless it is the only choice

Which fairness should we use?

- A universal design principle
for generality, impose the least constraint
no fairness < weak fairness < strong fairness
- In practice, fairness derives from scheduling
local process scheduler usually ensures that every
process has an opportunity to execute, if enabled
a process is enabled if any of its transitions is
enabled

Fair mutual exclusion, informally

- When process p_i wants to enter its critical region \mathcal{C} , it sets a flag f_i and stores i in shared variable $turn$
- While the flag of the other process is set and $turn$ has value i , p_i does not enter, but instead spins
- If the flag of the other process is not set or $turn$ has value other than i , p_i enters \mathcal{C}
- p_i resets f_i after leaving \mathcal{C}

Code skeleton (Peterson 2P)

shared variables: turn:1..2, initially undef,
read/write by all; f(i:1..2):Boolean, initially false,
writable by p_i , readable by all

process p_i :


f(i) := true

turn := i

while (f((i%2)+1) && turn = i)
 {null}

enter/use/exit critical region

f(i) := false



Safety property: mutual exclusion

- Need to argue that not possible for p_1 and p_2 to both be in critical region \mathcal{C} at same time

- Proof sketch

assume p_1 tries to enter \mathcal{C} , so f_1 set and $turn = 1$

if p_2 not competing, f_2 not set, so p_1 can enter \mathcal{C}

if p_2 is competing, f_2 is set, so entry depends on $turn$

assume p_1 is first, so p_2 sets $turn = 2$, which allows p_1 but blocks p_2

Liveness property: starvation free

- Need to argue that any process reaching T eventually enters C
- Proof sketch
 - if p_1 is waiting in T , it will be given priority through *turn* when p_2 exits from C
 - if p_2 competes, it will give priority to p_1 by setting *turn* to 2
- What level of fairness is supported?

Fair mutual exclusion, informally

- How can we generalize from 2 processes to n ?
- Use Peterson 2P iteratively in a series of $n-1$ competitions, each with its own variable *turn*
- At each competition level k , Peterson 2P ensures at least one loser i whose $turn_k$ is i if all compete
at level 1, at most $n-1$ processes proceed
at level 2, at most $n-2$ processes proceed
at level $n-1$, at most 1 process proceeds

Code skeleton (Peterson nP)

shared variables: $\text{turn}(k:1..n-1):1..n$, initially undef, read/write by all; $f(i:1..n):1..n-1$, initially 0, writable by p_i , readable by all

process p_i :

```
for k=1 to n-1 do {  
    f(i) := k  
    turn(k) := i  
    while (( $\exists j \neq i: f(j) \geq k$ ) && turn(k) = i) {null}  
}  
enter/use/exit critical region  
f(i) := 0
```

↑
↓ **T** (try)
↑
↓ **C** (critical)
↑
↓ **E** (exit)

Suitability for network environment

Q: What is a simple way to make these mutual exclusion algorithms suitable for use in a network environment?

i.e., how do we simulate shared memory?

A: Shared data can be encapsulated in processes that support communication via message passing rather than read/write